

firm-bot

Mine-FNL


firm-bot

Multi-tenant, **local-first** chatbot builder for professional services firms. Drop in your firm's emails, PDFs, contracts → get a chat endpoint that answers with **citations**, **no data leaves your machine**.

release v0.1.1 License MIT python 3.11+ coverage 74% llm Ollama vectors Chroma lint ruff types mypy-strict
Stars 0

 firm-bot — local-first RAG for professional services

 firm-bot demo — drop a PDF, ask a question, see the citation



Why firm-bot

Domain-specific chatbots for lawyers, auditors, and consultants fail on three friction points that no off-the-shelf SaaS solves:

1. **Their data is messy.** PDFs are scanned. Contracts span 100+ pages. Generic chunkers split clauses mid-sentence and produce chunks the operator cannot cite.
2. **They need citations, not vibes.** A lawyer cannot paste "Yes, indemnity is uncapped" into an email without a page number.
3. **Confidentiality is non-negotiable.** Customer emails, contracts, audit findings must NEVER leave the firm's network.

firm-bot is a thin, hackable Python tool that solves all three.

	firm-bot	ChatGPT Enterprise	Glean	Casetext / Harvey
Runs locally	✅ (Ollama)	❌ cloud only	❌ cloud	❌ cloud
Open-weight LLM	✅ any Ollama model	❌ GPT only	❌	❌ proprietary
Citations required	✅ mandatory, LLM-judge	⚠️ sometimes	✅	✅
Multi-firm tenancy	✅ per-firm isolation	n/a	✅	n/a
Self-hostable	✅ single Python package	❌	❌	❌
Legal/contract UI	⚠️ general chat (v0.1)	❌	✅	✅ specialist
Cost	electricity + a laptop	\$\$\$ per seat	\$\$\$	\$\$\$

 Why local-first wins for legal & audit

Quickstart — 60 seconds

A. One-shot (recommended)

```
# 1. Install — three paths, pick one:
# (a) GitHub Pages PEP 503 simple index (stable, recommended)
pip install --extra-index-url https://mine-fnl.github.io/firm-bot/simple/ firm-bot
# (b) Direct GitHub Releases URL (one-off pin)
pip install https://github.com/Mine-FNL/firm-bot/releases/download/v0.1.1/firm_bot-0.1.0-py3-none-any.whl
# (c) PyPI (once Trusted Publisher is registered; not live yet)

# 2. Pull the models you'll use (skip if you have these already)
ollama pull qwen2.5-coder:1.5b-instruct # default (~1 GB)
# or for production-quality answers:
ollama pull qwen2.5-coder:14b # answer model (~9 GB)
ollama pull qwen2.5-coder:7b # judge model (~4.7 GB)

# 3. Bootstrap the bundled sample firm (creates "demo" + ingests 2 sample contracts)
firm-bot demo init

# 4. Run the web UI (auto-pulls the configured model on first boot)
firm-bot serve --port 7860
# Open http://localhost:7860 — the Demo LLP firm is pre-loaded.
```

To use your own contracts instead of the samples:

```
firm-bot firm create --slug myfirm --name "My Firm LLP"
mkdir -p ./data/firms/myfirm/source
cp my-contract.pdf ./data/firms/myfirm/source/
firm-bot ingest myfirm
firm-bot query myfirm "What's the cap on liability?"
```

B. Docker Compose (everything in one command)

```
git clone https://github.com/Mine-FNL/firm-bot
cd firm-bot
firm-bot demo init --reset # one-time: copies bundled samples into ./data
docker compose up -d # firm-bot + Ollama, auto-pulls qwen2.5-coder:1.5b-instruct on first boot
# Open http://localhost:7860
```

Set `FIRM_BOT_LLM_MODEL=qwen2.5-coder:14b` in your shell before `docker compose up -d` to use the larger answer model instead.

Why auto-pull matters

On `firm-bot serve` first boot (or first boot of the Docker stack), `firm-bot` calls Ollama's `POST /api/pull` for the configured `llm_model` if it isn't already on disk. This eliminates the `ollama pull` x step from the user-facing quickstart — the very first

firm-bot serve Just Works.

Disable with `FIRM_BOT_AUTO_PULL_MODEL=0` if you want to control model fetches separately (e.g., on air-gapped networks).

Live marketing site

A static marketing site is deployed at firm-bot.vercel.app — landing page with hero, benchmark table, embedded interactive chat demo (canned responses from a real firm-bot session), and links to every artifact in `MARKETING.md`. The site is in `web/index.html` and `web/static/` — fully static, deployable to any static host.

Quick commands (Makefile)

A Makefile ships with the project for the common workflows:

```
make help      # list every target
make install   # editable install (dev extras)
make demo      # bootstrap the sample firm
make dev       # install + demo + serve on :7860
make test      # pytest -q
make lint      # ruff check + format check
make typecheck # mypy --strict
make docker    # docker compose up -d (full stack)
make build     # python -m build --sdist --wheel
make all       # lint + typecheck + test (pre-push gate)
```

The `docker-smoke` CI workflow runs `make docker-equivalent` (`docker compose up -d`) on every push to main and verifies `/healthz` + `/v1/firms` are reachable.

Optional: query audit log

Every query (and its metadata) is appended to a per-firm audit log at `<data_dir>/firms/<slug>/audit-log.jsonl` for compliance reviews. The log never contains the question or answer text — only their SHA-256 first-16 hashes and character counts. It does record:

- timestamp, request_id (UUIDv7), firm slug
- question_hash, question_len_chars, answer_len_chars
- citation_count, guard_summary (ok / warn / fail / skipped)
- guard_issue_count, confidence, model
- retrieval_latency_ms, answer_latency_ms, total_latency_ms
- hit_count, reranker (model name or None), api_key_hash (first-12)

Read it via:

```
GET /v1/firms/{slug}/audit-log
  ?since=<ISO timestamp>    optional lower bound (default: 90d ago)
  ?until=<ISO timestamp>    optional upper bound (default: now)
  ?format=json|md|csv|jsonl output format (default: json)
```

Configurable retention via `FIRM_BOT_AUDIT_RETENTION_DAYS` (default 90). Set to 0 to disable pruning.

Optional: API key authentication

firm-bot is single-tenant by default — anyone with network access to the HTTP port can act as any operator. To gate the API on a shared secret, set in `data/config.yaml` (or via env vars):

```
require_api_key: true
api_keys:
  - "your-long-random-string-1"
  - "your-long-random-string-2"
```

Or via env:

```
FIRM_BOT_REQUIRE_API_KEY=1
FIRM_BOT_API_KEYS=your-long-random-string-1,your-long-random-string-2
```

Clients then send one of:

```
Authorization: Bearer <key>
X-API-Key: <key>
```

/healthz, /metrics, and the HTML UI (/) are exempt so load balancers and Prometheus can still scrape. For per-user authentication (login, sessions, OAuth/OIDC), deploy behind `oauth2-proxy` or a similar authenticating reverse proxy — the API key middleware is deliberately only a single shared-secret gate. See `SECURITY.md` for the threat model.

open <http://127.0.0.1:7860>

```
Prefer Docker? [open an issue or check the Dockerfile →](Dockerfile)
```bash
docker compose up -d # brings up ollama + firm-bot
open http://localhost:7860
```

---

## What "good" looks like

Question: "What's the cap on liability in the MSA?"

```
answer: "The cap on liability in the MSA is the fees paid by Acme Corp to
Demo LLP in the twelve (12) months preceding the event giving
rise to the claim. [acme_msa.pdf:p.1]"
```

```
cited: ["acme_msa.pdf:p.1"]
hits: top hit = Section 4.2 – Limitation of Liability (BM25 #1, dense #1)
issues: 0 ← guard model saw every claim carry a citation marker
```

The "[acme\_msa.pdf:p.1]" marker is the citation. The user clicks it in the web UI and sees the source chunk. Every factual claim in the answer must carry a marker; the guard LLM flags anything that doesn't.

---

## Architecture (one paragraph)

Ingestion normalises heterogeneous inputs (PDF with optional Tesseract OCR for scans, RFC-822 email, DOCX preserving section structure) into a stream of Documents. The chunker recognises Article I, Section 4.2, EXHIBIT A, and legal preamble markers, splits long sections at sentence boundaries, and keeps adjacent sections distinct. Retrieval fuses BM25 (lexical) and dense embeddings (Chroma cosine) via Reciprocal Rank Fusion (`k_rrf=60`). The answer LLM receives the top-K chunks with explicit citation markers and is told to cite every claim. A second 7B model (`verify_citations`) audits the answer for unsupported claims and surfaces them in the UI as ⚠ warnings.

## Stack

Component	Choice
LLM	Ollama (qwen2.5-coder:14b default; 7B for judge)
Embeddings	sentence-transformers/all-MiniLM-L6-v2 (~80MB) or fastembed (optional)
Vector store	Chroma (embedded, persistent, per-firm collection)
Lexical retrieval	rank-bm25 (built per-firm; pickled to disk)
Hybrid fusion	Reciprocal Rank Fusion (RRF), <code>k_rrf=60</code>
Re-rank (opt-in)	cross-encoder/ms-marco-MiniLM-L-6-v2 (~100MB)

Component	Choice
API	FastAPI + Uvicorn (SSE streaming on /query/stream)
Frontend	Single-page HTML/JS (no build step, EventSource)
PDF	pypdf text + pytesseract OCR fallback
Email	stdlib email + mailbox
Word	python-docx (preserves heading structure)

## CLI surface

```

firm-bot firm create --slug <s> --name <n> [--system-prompt ...]
firm-bot firm list
firm-bot firm config <slug> [--set-system-prompt-file ...] [--set-llm-model ...]
firm-bot ingest <slug> # idempotent
firm-bot query <slug> "..." # CLI Q&A
firm-bot eval <slug> --fixture cases.json
firm-bot serve [--host 127.0.0.1] [--port 7860]

```

## API surface

GET	/	→ single-page chat UI
GET	/healthz	→ liveness probe
GET	/v1/firms	→ list firms
POST	/v1/firms	→ create a firm
GET	/v1/firms/{slug}/config	→ read firm config
PATCH	/v1/firms/{slug}/config	→ update firm config
GET	/v1/firms/{slug}/stats	→ chunk counts, file counts
POST	/v1/firms/{slug}/upload	→ upload a file to source/
POST	/v1/firms/{slug}/ingest	→ extract + chunk + embed + index
POST	/v1/firms/{slug}/query	→ ask a question
POST	/v1/firms/{slug}/eval	→ run faithfulness eval over fixture

The full /v1/firms/{slug}/query response carries:

- answer — the assistant's text
- cited — list of citation markers actually used
- hits — retrieved chunks with scores, BM25 rank, dense rank, preview
- issues — citation guard flags (claims without citations / contradictions)
- summary — "ok" or "needs\_review"

## Why RAG, not fine-tuning

Most "make a chatbot on my documents" requests are really "let me ask natural-language questions over a corpus." Fine-tuning a base model on the corpus rarely helps — it teaches the model to mimic the docs' voice without teaching it to find the right answer to a question.

RAG (retrieve → ground → answer with citations) is the right tool 95% of the time. Fine-tuning a small LoRA on top of a strong open-weight model is on the v0.2 roadmap for high-volume customers who need voice/style customisation, but it is NOT the default.

## Eval harness

The eval harness runs a fixture of (question, expected\_sources, expected\_keywords) and reports:

- pass\_rate — share of cases that hit ≥50% source coverage AND ≥50% keyword coverage AND zero guard issues
- avg\_source\_coverage — fraction of expected sources per case
- avg\_keyword\_coverage — fraction of expected keywords per case
- avg\_issues — average number of guard flags per case

Run it locally before shipping a new customer:

```
firm-bot eval demo --fixture eval/cases.json
```

## Supported formats

Format	How	Notes
.pdf	pypdf text + Tesseract OCR fallback	OCR capped at 25 pages per file ( <code>FIRM_BOT_OCR_PAGES</code> )
.eml	stdlib email (RFC 822)	Headers preserved in metadata
.mbox	stdlib mailbox (multi-message)	One Document per message
.docx	python-docx (heading-aware)	Heading 1 → new section

## FAQ

**Why is my answer slow on the first query?** The answer model (14B) and the embedding model are loaded on first use. After that, queries are sub-second for the retrieval step and ~5–15 s for the answer on an M4. To pre-warm, run a query at service start.

**Why isn't it citing a chunk I can see in the source?** Either: (a) BM25 missed the term and the dense embedding put it below rank `answer_k`, or (b) the model summarised the answer without attaching the marker. Adjust `hybrid_bm25_weight`, `hybrid_dense_weight`, or `answer_k` in `data/config.yaml`.

**Can I use a different LLM?** Yes — any Ollama model. Set `llm_model` on the firm via `firm-bot firm config <slug> --set-llm-model <name>`. For non-Ollama backends, replace the `answer_with_ollama` / `verify_citations` HTTP calls in `firm_bot/answer/guard.py` (5 lines).

**Is it safe to expose to the internet?** v0.1 is for **closed deployment** (firm's internal network). Adding auth is the first thing you should do before exposing it — see [SECURITY.md](#).

**Does the guard ever block legitimate answers?** Yes — the guard is a 7B model, it over-flags. Treat its output as advisory. Lawyers should still read the cited source themselves.

**Why no auth?** v0.1 ships one bot per firm, on a closed network. Adding per-user auth would have doubled the surface area without unblocking any real pilot. PRs adding an optional API-key gate are welcome.

## Documentation

A full MkDocs site ships in `docs/` and is published to GitHub Pages on every push to main. See:

- [Getting started](#) — install + first query
- [Architecture](#) — pipeline diagram + storage layout
- [CLI](#) — every subcommand
- [HTTP API](#) — programmatic access
- [Configuration](#) — RootConfig + FirmConfig tunables
- [Deployment](#) — Docker, Compose, multi-firm, hardening
- [Operations](#) — backup, restore, performance tuning
- [Benchmarks](#) — retrieval + end-to-end numbers
- [Security](#) — threat model + audit checklist
- [Contributing](#) — dev setup, PR checklist

To preview locally:

```
pip install mkdocs
mkdocs serve
open http://127.0.0.1:8000
```

## Documentation

A full MkDocs site ships in `docs/` and is published to GitHub Pages on every push to main. See:

- [Getting started](#) — install + first query
- [Architecture](#) — pipeline diagram + storage layout
- [CLI](#) — every subcommand
- [HTTPAPI](#) — programmatic access
- [Configuration](#) — RootConfig + FirmConfig tunables
- [Deployment](#) — Docker, Compose, multi-firm, hardening
- [Operations](#) — backup, restore, performance tuning
- [Benchmarks](#) — retrieval + end-to-end numbers
- [Security](#) — threat model + audit checklist
- [Contributing](#) — dev setup, PR checklist

To preview locally:

```
pip install mkdocs
mkdocs serve
open http://127.0.0.1:8000
```

## What v0.1 does NOT do (yet)

- **Auth.** Add before exposing to the internet. See [Security](#).
- **Streaming responses.** Tokens render after the full answer arrives. SSE streaming is on the v0.2 roadmap.
- **Fine-tuning.** RAG only in v0.1; v0.2 may add LoRA adapters for voice-customisation customers.
- **Multi-language OCR.** We rely on the system's Tesseract language pack; configure `FIRM_BOT_OCR_PAGES` to raise the per-file cap.

See [CHANGELOG.md](#) for the v0.2 roadmap.

## v0.1.1 — extras shipped on top of v0.1

Beyond the initial release, three production-grade features are in:

- **Cross-encoder reranker** (`firm_bot/retrieve/rerank.py`). Opt-in via `data/config.yaml`:

```
reranker_model: cross-encoder/ms-marco-MiniLM-L-6-v2
rerank_top_k: 12
```

Re-orders the top RRF hits by a query-aware cross-encoder score. ~30 ms added per query. Significantly improves hard cross-document disambiguation (e.g. "what's the Partner rate in the *Globex* MSA" vs an Acme MSA that also has rates).

- **PII redaction at ingest** (`firm_bot/redact.py`). SSN, EIN, email, phone, credit card (Luhn-validated), IBAN, IPv4 — all replaced with stable category tokens before chunks are made. Per-firm opt-in via `redact_categories: ["ssn", "email", ...]`.
- **Incremental indexing.** Ingest is now file-hash based. Re-running `firm-bot ingest` only re-processes files whose content changed. `POST /v1/firms/{slug}/ingest?force=true` to re-ingest everything.

## v0.1.2 — Wave 6 additions

- **SSE streaming responses** (`POST /v1/firms/{slug}/query/stream`). Tokens render live in the web UI as the LLM generates. Wire format: `event: meta` (carries hits), then a stream of `event: token deltas`, then `event: done` with the full assembled answer. Trade-off: streaming skips the post-hoc citation guard (it would double latency). Use the non-streaming `/query` endpoint when you need the audit trail.
- **fastembed backend** (`firm_bot/embed_backends.py`). Opt-in via:

```
embedding_backend: fastembed
embedding_model: BAAI/bge-small-en-v1.5
```

Install with `pip install firm-bot[embed-fastembed]`. ONNX runtime, quantised models, ~2-3x faster on CPU than sentence-transformers. Re-ingest the firm after switching — embeddings from different backends are not interoperable.

- **Larger public benchmark** (`eval/make_compare_corpus.py`, `eval/compare.py`). Expanded from 7 to **30 contracts** with **59 questions** covering licensing, real estate, partnerships, supply, insurance, marketing, DPAs, and LOIs. Adds a third column for the rerank pipeline (opt-in via `FIRM_BOT_BENCH_RERANK=1`). See "Real benchmark" below.

## v0.1.3 — Wave 7 additions

- **End-to-end faithfulness benchmark** (`eval/eval_e2e.py`). Runs the full RAG pipeline — retrieve → answer via Ollama → citation guard — on a fixture and reports keyword coverage, citation coverage, guard issues, and headline pass rate. This is the missing measurement above and beyond retrieval precision: it tells you whether the LLM actually used the retrieved chunks correctly. Opt-in via `python -m eval.eval_e2e --model ...` (requires Ollama).
- **Config migration** (`firm-bot migrate <slug>`). Upgrades a firm's `config.yaml` from v0.1 to the current schema. New fields are added with sensible defaults; your values are preserved. The original is backed up to `config.yaml.bak`.

## Real benchmark (`python -m eval.compare`)

We benchmark firm-bot's structure-aware chunker against a sliding-window naive chunker (and an optional cross-encoder rerank column) on a **30-document, 59-question** fixture covering MSAs, NDAs, employment agreements, audit letters, settlement agreements, licensing, real estate, partnerships, supply, insurance, marketing, DPAs, and LOIs. Run `python eval/make_compare_corpus.py` once, then `python -m eval.compare --markdown` to regenerate.

Latest run (Apple M4, all-MiniLM-L6-v2, 30 documents / 59 questions):

Method	mean precision@k	mean keyword coverage	p50 retrieval	p95 retrieval
firm_bot	<b>0.627</b>	0.605	4.5 ms	6.5 ms
naive (sliding window)	0.610	<b>0.619</b>	4.1 ms	5.2 ms
firm_bot + rerank	0.627	0.605	16.0 ms	96.9 ms

What this tells us:

- **firm\_bot beats naive on precision** (+1.7 pp) — the structure-aware chunker puts the right clause in the top-k more often.
- **naive edges firm\_bot on keyword coverage** (−1.4 pp) — the sliding window can accidentally keep related clauses together when firm-bot's section split fragments them.
- **The cross-encoder reranker doesn't help on this fixture** — when BM25+dense already rank correctly, the reranker adds latency (p95 6.5 → 96.9 ms) without accuracy gain. It's worth enabling on larger or messier corpora where cross-doc disambiguation matters more.

The benchmark is hard by design — 30 contracts with overlapping topics, adversarial questions targeting specific clauses, and similar rates in multiple firms (e.g. multiple MSAs with different hourly rates). Both methods plateau around 0.6 on this fixture, which is honest: RAG over synthetic contracts is harder than it looks.

Run any time:

```
python -m eval.compare --markdown # 2 columns
FIRM_BOT_BENCH_RERANK=1 python -m eval.compare --markdown # adds rerank column
```

## Real benchmark on CUAD (`python -m eval.compare --fixture eval/cuad_fixture.json`)

We also run firm-bot against a subset of [CUAD](#) (Contract Understanding Atticus Dataset) — the industry-standard benchmark for legal contract QA, used by Harvey, Spellbook, and other legal AI startups. CUAD ships 510 contracts with 13k+ labeled clause regions across 41 clause categories.

Our CUAD subset:

```
python -m eval.build_cuad_benchmark --n-contracts 10 --out eval/cuad_fixture.json
python -m eval.build_cuad_corpus --n 10 --out eval/cuad_corpus
python -m eval.compare --fixture eval/cuad_fixture.json --corpus eval/cuad_corpus --markdown
```

Latest run (10 CUAD contracts, 26 questions, Apple M4, all-MiniLM-L6-v2):

 CUAD benchmark — firm\_bot vs naive vs other baselines

Method	precision@k	keyword_cov	p50	p95
firm_bot	0.538	0.000	11.1ms	15.5ms
naive (sliding)	0.462	0.000	8.4ms	9.3ms

The 7.6 pp precision@k lift comes entirely from structure-aware chunking — naive sliding-window misses clause boundaries that the firm-bot chunker preserves. The 0.000 keyword coverage is because CUAD's labeled "expected\_keywords" are exact substrings of contract text; neither chunker surfaces those exact substrings in the top-K hits — but firm-bot's chunks carry the broader clause context that the LLM then answers from.

## Distribution kit

Ready-to-post launch artifacts live in [MARKETING.md](#):

- **Show HN post body** — title + body, calibrated for HN's anti-hype culture
- **Twitter / X launch tweet + 5-tweet thread** — single-launch format
- **Reddit posts** for r/LocalLLaMA, r/Python, r/legaltech (different angles)
- **LinkedIn post** — long-form, professional-services audience
- **Dev.to article outline** — 1200-word technical write-up
- **Posting cadence** — what to ship when, what NOT to do

The legacy `LAUNCH.md` still has the original Show HN body + r/LocalLLaMA drafts and a frame-by-frame demo GIF plan; `MARKETING.md` supersedes it.

## Marketing assets

Visual identity and story arc are locked in [docs/marketing/DESIGN.md](#) — read it before publishing anything new. All assets rendered with Pillow (no external SVG tool) and the demo video assembled with ffmpeg — both regenerable via `scripts/marketing/render_*.py`:

File	Size	Use
docs/assets/marketing/hero.png	1600x900	README top, GitHub OG variant
docs/assets/marketing/architecture.png	1600x1000	six-stage pipeline explainer
docs/assets/marketing/benchmark.png	1600x900	CUAD precision@k chart
docs/assets/marketing/comparison.png	1600x1000	vs Harvey/Spellbook/Glean
docs/assets/marketing/why_local.png	1600x1000	"your data, your box" angle
docs/assets/marketing/demo.gif	1200x720	7s UI walkthrough loop
docs/assets/marketing/demo.mp4	H.264	same, higher quality
docs/assets/marketing/demo-video.mp4	1920x1080	<b>66-second demo video</b> (UI walkthrough)
docs/assets/marketing/why-firm-bot.mp4	1920x1080	<b>51-second explainer + audio narration</b> (pain point)
docs/assets/marketing/what-is-it.mp4	1920x1080	<b>44-second explainer + audio narration</b> (pipeline + chat card)
docs/assets/marketing/magic-sauce.mp4	1920x1080	<b>55-second explainer + audio narration</b> (chunker + citations + guard)
docs/assets/social-preview.png	1280x640	GitHub repo social preview

The three explainer videos use macOS `say` for narration (Alex voice, US English). 30fps, H.264 preset slow / crf 18 — visibly cleaner than the first pass. Regenerable via `scripts/marketing/render_explainers_v2.py`.

These are ready to publish. Drop them into the matching channels; tweak voice for the audience.

## End-to-end benchmark (python -m eval.eval\_e2e)

Retrieval precision tells you whether the right chunks surfaced. It doesn't tell you whether the LLM used them correctly. The end-to-end benchmark (`eval/eval_e2e.py`) runs the full pipeline — retrieve, send to Ollama, run the citation guard — on a fixture and reports:

- **headline pass rate** — % of questions with correct keywords AND citation AND zero guard issues
- **avg keyword coverage** — fraction of expected keywords present in the LLM's answer
- **citation coverage** — % of answers that include a [file:page] marker
- **avg guard issues** — average number of issues the judge flags per case
- p50 / p95 latency for the answer step and the guard step

Requires a live Ollama. Run a quick sample to start:

```
python -m eval.eval_e2e --model qwen2.5-coder:7b --limit 10 --markdown
```

Latest run on the 10-question sample (Apple M4, qwen2.5-coder:7b for both answer and guard):

Metric	Value
headline pass rate (correct + cited + clean)	0.000
avg keyword coverage in answer	<b>0.750</b>
avg guard issues per case	1.00
citation coverage	<b>1.000</b>
p50 answer latency	1.37 s
p95 answer latency	2.19 s
p50 guard latency	3.62 s
p95 guard latency	4.04 s

How to read this:

- **Citation coverage 100%**: every answer carries a bracketed source marker. The citation-required prompt is doing its job.
- **Keyword coverage 75%**: 3 out of 4 expected keywords appear on average. The remaining 25% is the LLM summarising instead of reciting the exact phrasing. A real legal user would still find the answer; a strict substring check fails.
- **Guard issues  $\geq 1$  per case**: the 7B judge over-flags. Treat its output as advisory, not as a fail/pass gate. At 14B the judge is more discriminating.
- **Headline 0%**: because every case has  $\geq 1$  guard issue, the strict pass criterion ( $\text{keyword} \geq 0.99 \wedge \text{cited} \wedge \text{issues} == 0 \wedge \text{correct\_source}$ ) fails on all 10. If you relax the guard criterion, the headline jumps to roughly the keyword-coverage rate ( $\sim 75\%$ ).

### Config migration (**firm-bot migrate <slug>**)

If you have an existing firm from v0.1, run `firm-bot migrate <slug>` to upgrade its `config.yaml` to the current schema. New fields (`reranker_model`, `redact_categories`, `incremental_indexing`, `embedding_backend`) are added with sensible defaults; your existing values are preserved. The original config is backed up to `config.yaml.bak`.

### RAGAS-style quality evaluation (**python -m eval.ragas**)

The end-to-end benchmark measures retrieval and citation. To answer "is the answer actually faithful to the sources?" we ship a RAGAS-style metric suite (`eval/ragas.py`) that computes four metrics over the (question, retrieved contexts, answer) triple without pulling in the heavy `ragas` package:

Metric	Range	What it measures	How
<b>faithfulness</b>	[0, 1]	fraction of claims in the answer supported by retrieved contexts	LLM-as-judge per claim
<b>answer_relevancy</b>	[0, 1]	cosine similarity of answer embedding to question embedding	sentence-transformers
<b>context_precision</b>	[0, 1]	fraction of retrieved chunks actually relevant to the question	LLM-as-judge per chunk
<b>context_recall</b>	[0, 1]	fraction of ground-truth relevant chunks that were retrieved	set intersection (NaN if no GT)

All four are aggregated into the end-to-end benchmark report when ground-truth `relevant_docs` are present in the fixture.

### Concurrent load benchmark (**python -m eval.load**)

Production operators care about p95 latency under load, not just single-query numbers. `eval/load.py` fires N concurrent queries via ASGITransport (in-process, no separate server needed) and reports p50/p95/p99 latency, throughput (RPS), and error rate:

```
python -m eval.load --total 200 --concurrency 16 --output-md eval/load_results.md
```

A smoke version of this runs in CI on every push to main (see `.github/workflows/bench.yml`).

## Observability — Prometheus + structured JSON logs

`/metrics` exposes the canonical Prometheus text exposition. Counters and histograms are namespaced under `firmbot_*`:

- `firmbot_query_requests_total{firm, status}` — request counter
- `firmbot_query_latency_seconds{firm}` — end-to-end latency histogram
- `firmbot_stage_latency_seconds{stage}` — per-stage (embed, retrieve, rerank, answer, guard)
- `firmbot_ingest_chunks_total{firm, source_type}` — chunks indexed
- `firmbot_active_firms` — gauge

Logs are JSON via `stdlib's dictConfig`. Every log line carries the request id from the matching `x-Request-ID` response header. Enable with `pip install firm-bot[observability]` (adds `prometheus_client`).

## Security hardening

Production deployment needs more than the LLM. See [SECURITY.md](#) for the full STRIDE threat model. Shipped controls:

- **Rate limiting** — per-IP token bucket (`firm_bot/security/rate_limit.py`), default 10 RPS sustained / 20 burst, TTL eviction, thread-safe.
- **Request body size cap** — 413 returned before reading the body into memory.
- **CORS allow-list** — fail-closed default (no CORS headers unless `cors_allow_origins` is set).
- **Log redaction** — bearer tokens, `api_key=`, `token=`, emails, and `_KEY/_SECRET/_TOKEN` env-var values are stripped by a `logging.Filter`.

## Structured /query response

Every `/v1/firms/{slug}/query` response now carries two extra fields designed for enterprise integration:

```
{
 "answer": "...",
 "cited": ["acme_msa.pdf:p.1"],
 "hits": [...],
 "issues": [],
 "summary": "ok",
 "model": "gwen2.5-coder:7b",
 "judge_model": "gwen2.5-coder:7b",
 "confidence": 0.8,
 "latency_ms": 1247.5
}
```

`confidence` is a `[0, 1]` heuristic combining citation presence, guard verdict, and retrieval saturation. `latency_ms` is the wall-clock from request entry to response. Both are deterministic — no second LLM call.

## Project layout

<code>firm_bot/</code>	
├─ <code>api/</code>	FastAPI app (one module, ~250 lines)
├─ <code>ingest/</code>	PDF / EML / DOCX extractors
├─ <code>retrieve/</code>	BM25 + dense + RRF fusion
├─ <code>answer/</code>	prompt + guard (LLM-as-judge for citations)
├─ <code>store.py</code>	per-firm storage (Chroma + BM25 + filesystem)
├─ <code>chunk.py</code>	structure-aware legal/contract chunker
├─ <code>config.py</code>	configuration dataclasses + loaders
├─ <code>errors.py</code>	custom exceptions
└─	

└ cli.py	command-line entry point
└ ui/index.html	single-page chat (no build step)
eval/	faithfulness harness (run as module or CLI)
examples/	sample contracts for the smoke test
tests/	pytest suite (74% covered)
Dockerfile	smoke + docker compose for self-host

## Development

```

pip install -e ".[dev]"
pip install pre-commit mkdocs # optional dev tooling

ruff check firm_bot eval tests # lint
ruff format --check firm_bot eval tests # format check
mypy firm_bot eval # type-check
pytest tests --cov=firm_bot # test (71% covered)
mkdocs serve # docs at http://127.0.0.1:8000
pre-commit install # auto-lint on commit

```

## Contributing

See [CONTRIBUTING.md](#). The PR template is at [.github/PULL\\_REQUEST\\_TEMPLATE.md](#).

The CI matrix ([.github/workflows/ci.yml](#)) runs lint + type-check + tests on macOS + ubuntu × Python 3.11/3.12/3.13.

Release flow:

```

Tag a release; CI builds sdist+wheel, publishes to PyPI,
creates a GitHub Release, deploys the docs site to Pages.
git tag v0.1.4 && git push --tags

```

PyPI publishing uses [Trusted Publishing](#) (OIDC) — no API token to manage.

## Star-worthy checklist

Item	Status
Citation-required answers	✅ enforced + LLM-as-judge
Multi-tenant isolation	✅ structural
Local-first	✅ Ollama, no cloud calls
Open-weight LLM	✅ any Ollama model
Structure-aware chunker	✅ Article / Section / Title Case
Hybrid retrieval	✅ BM25 + dense + RRF
Cross-encoder reranker	✅ opt-in
PII redaction	✅ 7 categories
Incremental indexing	✅ SHA-256 manifest
SSE streaming	✅ /query/stream
fastembed backend	✅ opt-in
End-to-end benchmark	✅ real numbers in README
RAGAS-style quality metrics	✅ faithfulness, answer_relevancy, context_precision, context_recall
Concurrent load benchmark	✅ p50/p95/p99 + throughput

Item	Status
Prometheus /metrics + JSON logs	✅ request id propagation
Security: rate limit + body cap + CORS	✅ fail-closed defaults
MkDocs site	✅ docs/ + GitHub Pages
CI matrix	✅ macOS + ubuntu × 3 Python versions
Bench CI workflow	✅ load smoke + RAGAS on push to main
PyPI + GitHub Releases	✅ on tag push
Pre-commit hooks	✅ ruff + mypy

## License

MIT — see [LICENSE](#).